

07-2 Device Driver Basics

Using kernel modules

Free Electrons

Loadable kernel modules

- ▶ Modules: add a given functionality to the kernel (drivers, filesystem support, and many others)
- ▶ Can be loaded and unloaded at any time, only when their functionality is need
- ▶ Useful to keep the kernel image size to the minimum (essential in GNU/Linux distributions for PCs)
- ▶ Also useful to reduce boot time: you don't spent time initializing devices and kernel features that you only need later
- ▶ Caution: once loaded, have full access to the whole kernel address space. No particular protection

Minimal Device Driver (Listing 8-1)

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>
static int __init hello_init(void) {
    printk(KERN_INFO "Hello Example Init\n");
    return 0;
}
static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}
module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

Module Build Infrastructure

1. Starting from the top-level Linux source directory, create a directory under `.../drivers/char` called **examples**.
2. Add a menu item to the kernel configuration to enable building **examples** and to specify built-in or loadable kernel module. See Section 8.1.4, page 205
3. Add the new examples subdirectory to the `.../drivers/char/Makefile` conditional on the menu item created in step 2.
4. Create a **Makefile** for the new **examples** directory, and add the **hello1.o** module object to be compiled conditional on the menu item created in step 2.
5. Finally, create the driver **hello1.c** source file from Listing 8-1.

Typo page 206

```
diff --git a/drivers/char/Kconfig b/drivers/char/Kconfig
index 6f31c94..0805290 100644
--- a/drivers/char/Kconfig
+++ b/drivers/char/Kconfig
@@ -4,6 +4,13 @@

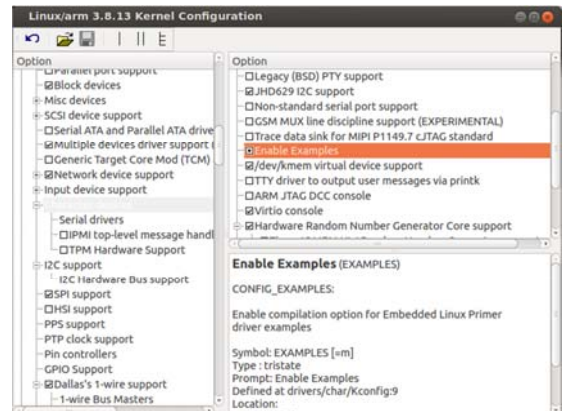
menu "Character devices"

+config EXAMPLES
+    tristate "Enable Examples"
+    default m
+    ---help---
+    Enable compilation option for Embedded Linux Primer
+    driver examples
+
config DEVKMEM
    bool "/dev/kmem virtual device support"
    default y
    help
```

Must be lower case

In new kernel

Check Config



Module Build Output

```
host$ time make modules
CHK    include/generated/uapi/linux/version.h
CHK    include/generated/utsrelease.h
make[1]: `include/generated/mach-types.h' is up to date.
CALL   scripts/checksyscalls.sh
CC [M]  drivers/char/examples/hello1.o
Building modules, stage 2.
MODPOST 1326 modules
LD [M]  drivers/char/examples/hello1.ko

real    0m33.706s
user    0m31.462s
sys     0m4.824s
```

First time takes much longer
Try make -jX modules

Once built... Option 1

On the Beagle... Two choices....

- Option 1:
`make INSTALL_MOD_PATH=~/.BeagleBoard modules_install`
- Will create **lib** directory in `~/BeagleBoard` with everything that goes in `/lib` on the Beagle

```
host$ ls -F ~/.BeagleBoard/lib/modules/3.8.13/+
```

```
build#          modules.dep.bin  modules.seriomap
kernel/         modules.devname  modules.softdep
modules.alias   modules.ieee1394map  modules.symbols
modules.alias.bin  modules.inputmap  modules.symbols.bin
modules.builtin  modules.isapnpmap  modules.usbmap
modules.builtin.bin  modules.ofmap  source#
modules.cwmmap   modules.order
modules.dep      modules.pcimap
```

- Then
`host$ rm build source`
`host$ scp -r ~/.BeagleBoard/lib root@beagle:/lib`
- Could take a while to transfer

Once built... Option 2

- Just copy the new file you created

```
host$ scp .../drivers/char/examples/hello1.ko root@beagle:.
```

- On the Beagle

```
beagle$ cd /lib/modules/3.8.13/kernel/drivers/char/
beagle$ mkdir examples
beagle$ cd examples
beagle$ mv ~/hello1.ko .
• Now build a new dependencies file
beagle$ depmod -a
```

Loading and Unloading a Module

```
beagle$ modprob hello1
beagle$ dmesg | tail -4
[  9.106206] snd-usb-audio 1-1:1.0: usb_probe_interface
[  9.106244] snd-usb-audio 1-1:1.0: usb_probe_interface - got id
[  9.813239] usbcore: registered new interface driver snd-usb-
audio
[ 109.308551] Hello Example Init
beagle$ rmmod hello1
beagle$ dmesg | tail -4
[  9.106244] snd-usb-audio 1-1:1.0: usb_probe_interface - got id
[  9.813239] usbcore: registered new interface driver snd-usb-
audio
[ 109.308551] Hello Example Init
[ 241.037368] Hello Example Exit
```

Module Utilities

```
$ insmod /lib/modules/`uname -r`/kernel/drivers/char/examples/hello1.ko
```

- No need build dependencies file

Example Driver with Parameter

```
/* Example Minimal Character Device Driver */
#include <linux/module.h>

static int debug_enable = 0;
module_param(debug_enable, int, 0);

MODULE_PARM_DESC(debug_enable, "Enable module debug mode.");

static int __init hello_init(void) {
    /* Now print value of new module parameter */
    printk("Hello Example Init - debug mode is %s\n",
           debug_enable ? "enabled" : "disabled");
    return 0;
}

static void __exit hello_exit(void) {
    printk("Hello Example Exit\n");
}

module_init(hello_init);
module_exit(hello_exit);

MODULE_AUTHOR("Chris Hallinan");
MODULE_DESCRIPTION("Hello World Example");
MODULE_LICENSE("GPL");
```

Passing Parameters to a Module

```
beagle$ modprobe hello2 debug_enable=1
Hello Example Init - debug mode is enabled
beagle$ rmmod hello2
beagle$ modprobe hello2
Hello Example Init - debug mode is disabled
```

Other module commands

```
# /sbin/lsmmod
# /sbin/modinfo hello1
# /sbin/rmmod hello1
# /sbin/depmod (creates modules.dep.bin)
```

- Go play with them

Adding File System Ops to Hello.c

- Section 8.3, page 217 has a long example about adding file system operations to **hello.c**
- Look it over
- Creates a new device (**/dev/hello1**)
- You can read and write it
- Do it

Driver File System Operations

- Once a device driver is loaded into the live kernel...
 - **open()** is used to prepare it for subsequent operations
 - **release()** is used to clean up
 - **ioctl()** is used for nonstandard communication
- Think in terms of reading and writing a file...

```
fd = open("file", ...
read(fd, ...
close(fd)
```

open/release additions to hello.c

```
#include <linux/fs.h>

#define HELLO_MAJOR 234
...
struct file_operations hello_fops;

static int hello_open(struct inode *inode, struct file *file) {
    printk("hello_open: successful\n");
    return 0;
}

static int hello_release(struct inode *inode, struct file *file) {
    printk("hello_release: successful\n");
    return 0;
}
```

read/write additions to hello.c

```
static ssize_t hello_read(struct file *file,
    char *buf, size_t count, loff_t *ptr) {
    printk("hello_read: returning zero bytes\n");
    return 0;
}

static ssize_t hello_write(struct file *file,
    const char *buf, size_t count, loff_t *ppos)
{
    printk("hello_read: accepting zero bytes\n");
    return 0;
}
```

ioctl additions to hello.c

```
static long hello_ioctl(struct inode *inode,
    struct file *file, unsigned int cmd,
    unsigned long arg) {
    printk("hello_ioctl: cmd=%ld, arg=%ld\n",
        cmd, arg);
    return 0;
}
```

init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
        debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello1", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");

    /* Init processing here... */

    return 0;

hello_fail1:
    return ret;
}
```

Major number for device driver

- Every device has a major and minor number

```
$ ls -ls /dev/console
0 crw----- 1 yoder root 5, 1 2011-02-06 17:57 /dev/console
```

- Device numbers *used* to be statically assigned
- See .../Documentation/devices.txt

```
5 char  Alternate TTY devices
          0 = /dev/tty           Current TTY device
          1 = /dev/console System console
          2 = /dev/ptmx          PTY master multiplex
          64 = /dev/cua0         Callout device for ttyS0
```

- The text uses static assignment

```
234-239          UNASSIGNED
240-254 char     LOCAL/EXPERIMENTAL USE
```

Registering our functions

- Struct file_operations is used bind our functions to the requests from the file system.

```
struct file_operations hello_fops = {
    owner:    THIS_MODULE,
    read:     hello_read,
    write:    hello_write,
    ioctl_ioctl: hello_ioctl,
    open:     hello_open,
    release:  hello_release,
};
```

init additions to hello.c

```
#define HELLO_MAJOR 234
static int __init hello_init(void)
{
    int ret;
    printk("Hello Example Init - debug mode is %s\n",
        debug_enable ? "enabled" : "disabled");
    ret = register_chrdev(HELLO_MAJOR, "hello1", &hello_fops);
    if (ret < 0) {
        printk("Error registering hello device\n");
        goto hello_fail1;
    }
    printk("Hello: registered module successfully!\n");
    /* Init processing here... */
    return 0;
hello_fail1:
    return ret;
}
```

Device Nodes and mknod

- Use **mknod** to create a new device

```
$ mknod /dev/hello1 c 234 0
```



- Then

```
$ ls -l /dev/hello1
crw-r--r-- 1 root root 234, 0 Apr 2 2011 /dev/hello1
```

Dynamic Major Number

- The above example uses the older *static* method to assign a device number
- Today dynamic allocation is preferred
- Here is how:

```
#include <linux/kdev_t.h>
dev_t dev;
```

- This declares **dev** to be a device number (both major and minor). Now assign it a value
- **dev = MKDEV(234, 0);**

Requesting a number

- Now request a number

```
#include <linux/fs.h>
```

```
int register_chrdev_region(dev, 4, "hello");
```

- This requests a device number starting with 234 (previous page)
- It asks for 4 minor numbers
- Uses the name "hello"

- When done with the device use:

```
void unregister_chrdev_region(dev, 4);
```

Using **mknod**

- If your major number is assigned dynamically, how do you use **mknod**? Try the following

```
module="hello"
device="hello"
mode="664"
# remove stale nodes
/sbin/insmod ./module.ko $* || exit 1
rm -f /dev/${device}0
major=$(awk "\$2==\"$module\" {print \$1} /proc/devices"
mknod /dev/${device}0 c $major 0
```

/proc/devices

Character devices:	89 i2c	Block devices:	70 sd
1 mem	90 mtd	1 ramdisk	71 sd
4 /dev/vc/0	116 alsa	259 blkext	128 sd
4 tty	128 ptm	7 loop	129 sd
4 ttys	136 pts	8 sd	130 sd
5 /dev/tty	153 spi	31 mtdblock	131 sd
5 /dev/console	161 ircomm	65 sd	132 sd
5 /dev/ptmx	166 ttyACM	66 sd	133 sd
7 vcs	180 usb	67 sd	134 sd
10 misc	189 usb_device	68 sd	135 sd
13 input	212 DVB	69 sd	179 mmc
14 sound	216 rfcomm		254 device-mapper
29 fb	226 drm		
81 video4linux	234 hello3		
	244 ttyGS		
	245 ttySDIO		
	246 usbmon		
	247 uio		
	248 ttyO		
	249 bsg		

Module dependencies

- ▶ Some kernel modules can depend on other modules, which need to be loaded first
- ▶ Example: the [usb-storage](#) module depends on the [scsi_mod](#), [libusual](#) and [usbcore](#) modules
- ▶ Dependencies are described in [/lib/modules/<kernel-version>/modules.dep](#)

/lib/modules/3.8.13/modules.dep

```
kernel/lib/raid6/raid6_pq.ko:
kernel/lib/ts_kmp.ko:
kernel/lib/ts_bm.ko:
kernel/lib/ts_fsm.ko:
kernel/lib/notifier-error-inject.ko:
kernel/lib/pm-notifier-error-inject.ko:
kernel/lib/notifier-error-inject.ko:
kernel/lib/mpi/mpi.ko:
kernel/lib/asn1_decoder.ko:
kernel/lib/oid_registry.ko:
kernel/drivers/char/examples/hello2.ko:
kernel/drivers/char/examples/hello3.ko:
kernel/drivers/char/examples/hello1.ko:
```

Kernel log

When a new module is loaded, related information is available in the kernel log

- ▶ The kernel keeps its messages in a circular buffer (so that it doesn't consume more memory with many messages)
- ▶ Kernel log messages are available through the `dmesg` command ("diagnostic message")
- ▶ Kernel log messages are also displayed in the system console (messages can be filtered by level using `/proc/sys/kernel/printk`)

printk

- `/proc/sys/kernel/printk`
- The four values in this file are
 - `console_loglevel`,
 - `default_message_loglevel`,
 - `minimum_console_level` and
 - `default_console_loglevel`.
- These values influence `printk()` behavior when printing or logging error messages
- Messages with a higher priority than `console_loglevel` will be printed to the console
- Messages without an explicit priority will be printed with priority `default_message_level`

<http://www.tin.org/bin/man.cgi?section=5&topic=proc>

Kernel log levels

0 (KERN_EMERG)	The system is unusable
1 (KERN_ALERT)	Actions that must be taken care of immediately
2 (KERN_CRIT)	Critical conditions
3 (KERN_ERR)	Noncritical error conditions
4 (KERN_WARNING)	Warning conditions that should be taken care of
5 (KERN_NOTICE)	Normal, but significant events
6 (KERN_INFO)	Informational messages that require no action
7 (KERN_DEBUG)	Kernel debugging messages, output by the

Module utilities (1)

- ▶ `modinfo <module_name>`
`modinfo <module_path>.ko`
Gets information about a module: parameters, license, description and dependencies.
Very useful before deciding to load a module or not.
- ▶ `sudo insmod <module_path>.ko`
Tries to load the given module. The full path to the module object file must be given.

Understanding module loading

- ▶ When loading a module fails, `insmod` often doesn't give you enough details!
- ▶ Details are often available in the kernel log
- ▶ Example:

```
beagle$ sudo insmod ./intr_monitor.ko
insmod: error inserting './intr_monitor.ko': -1
Device or resource busy
beagle$ dmesg
[17549774.552000] Failed to register handler for
irq channel 2
```

Module utilities (2)

- ▶ `sudo modprobe <module_name>`
Most common usage of `modprobe`: tries to load all the modules the given module depends on, and then this module. Lots of other options are available. `modprobe` automatically looks in `/lib/modules/<version>/` for the object file corresponding to the given module name.
- ▶ `lsmod`
Displays the list of loaded modules
Compare its output with the contents of `/proc/modules!`

lsmod

```
beagle$ lsmod
Module                  Size  Used by
hello3                  1952  0
iptable_nat             2493  0
nf_conntrack_ipv4      13161  1
nf_defrag_ipv4         1378  1 nf_conntrack_ipv4
nf_nat_ipv4             3487  1 iptable_nat
nf_nat                  16075  2 nf_nat_ipv4,iptable_nat
nf_conntrack            80162  4 nf_nat,nf_nat_ipv4,iptable_nat,nf_conntrack_ipv4
ip_tables               11013  1 iptable_nat
x_tables                16948  1 ip_tables
g_multi                 56699  2
libccomposite          17303  1 g_multi
rfcomm                  30945  0
ircomm_tty             17315  0
ircomm                  10198  1 ircomm_tty
irda                    108219  2 ircomm_tty,ircomm
hidp                    12626  0
bluetooth               190961  4 hidp,rfcomm
..
```

Module utilities (3)

- ▶ **sudo rmmod <module_name>**
Tries to remove the given module.
Will only be allowed if the module is no longer in use
(for example, no more processes opening a device file)
- ▶ **sudo modprobe -r <module_name>**
Tries to remove the given module and all dependent modules (which are no longer needed after the module removal)

Passing parameters to modules

- ▶ Find available parameters:
`modinfo snd-intel8x0m`
- ▶ Through insmod:
`sudo insmod ./snd-intel8x0m.ko index=-2`
- ▶ Through modprobe:
Set parameters in `/etc/modprobe.conf` or in any file in `/etc/modprobe.d/`:
`options snd-intel8x0m index=-2`
- ▶ Through the kernel command line,
when the module is built statically into the kernel:

`snd-intel8x0m.index=-2`

module name ↑
module parameter name ↑
module parameter value —

Useful reading

Linux Kernel in a Nutshell, Dec 2006

- ▶ By Greg Kroah-Hartman, O'Reilly
<http://www.kroah.com/lkn/>
- ▶ A good reference book and guide on configuring, compiling and managing the Linux kernel sources.
- ▶ **Freely available on-line!**
Great companion to the printed book
for easy electronic searches!
Available as single PDF file on
<http://free-electrons.com/community/kernel/lkn/>
- ▶ In **exercises/pptx**



Useful reading too

Linux Device Drivers, Third Edition, February 2005

- ▶ By Jonathan Corbet, Alessandro Rubini, Greg Kroah-Hartman, O'Reilly
<http://lwn.net/Kernel/LDD3/>
- ▶ **Freely available on-line!**
Great companion to the printed book
for easy electronic searches!
Available as single PDF file
- ▶ LDD3 is current as of the 2.6.10 kernel (Old?)
- ▶ In **exercises/pptx**

